

SDN DEMO PREP

Iteration 2

By

Kabiru Sanusi

Hands-on Training with Python Net Apps - Elective 1

```
File Machine View Input Devices Help

ks-sdn@sdn_server:~$ ps -ea | grep ov
1331 ?      00:00:00 ovsdb-server
1456 ?      00:00:00 ovs-vswitchd
1602 ?      00:00:00 ovs-testcontrol
ks-sdn@sdn_server:~$ sudo ovs-vsctl show
[sudo] password for ks-sdn:
15edd720-f5b6-40f4-ab15-0f05b3615fc1
    ovs_version: "2.9.8"
ks-sdn@sdn_server:~$ sudo ovs-vsctl add-br br0
ks-sdn@sdn_server:~$ sudo ovs-vsctl add-port br0 veth0
ovs-vsctl: Error detected while setting up 'veth0': could not open network device veth0 (No such device). See ovs-vswitchd log for details.
ovs-vsctl: The default log directory is "/var/log/openvswitch".
ks-sdn@sdn_server:~$ sudo ovs-vsctl add-port br0 vif1.0
ovs-vsctl: Error detected while setting up 'vif1.0': could not open network device vif1.0 (No such device). See ovs-vswitchd log for details.
ovs-vsctl: The default log directory is "/var/log/openvswitch".
ks-sdn@sdn_server:~$ sudo ovs-vsctl show
15edd720-f5b6-40f4-ab15-0f05b3615fc1
    Bridge "br0"
        Port "br0"
            Interface "br0"
                type: internal
        Port "veth0"
            Interface "veth0"
                error: "could not open network device veth0 (No such device)"
        Port "vif1.0"
            Interface "vif1.0"
                error: "could not open network device vif1.0 (No such device)"
    ovs_version: "2.9.8"
ks-sdn@sdn_server:~$
ks-sdn@sdn_server:~$ sudo ovs-vsctl -- --if-exists del-port br0 veth0
ks-sdn@sdn_server:~$ sudo ovs-vsctl -- --if-exists del-br br0
ks-sdn@sdn_server:~$ sudo ovs-vsctl show
15edd720-f5b6-40f4-ab15-0f05b3615fc1
    ovs_version: "2.9.8"
ks-sdn@sdn_server:~$
```

File Machine View Input Devices Help

```
ks-sdn@sdn_server:~$ sudo mn --topo=single,3 --controller=none --mac
[sudo] password for ks-sdn:
*** Creating network
*** Adding controller
*** Adding hosts:
h1 h2 h3
*** Adding switches:
s1
*** Adding links:
(h1, s1) (h2, s1) (h3, s1)
*** Configuring hosts
h1 h2 h3
*** Starting controller

*** Starting 1 switches
s1 ...
*** Starting CLI:
mininet> dump
<Host h1: h1-eth0:10.0.0.1 pid=13374>
<Host h2: h2-eth0:10.0.0.2 pid=13376>
<Host h3: h3-eth0:10.0.0.3 pid=13378>
<OVSSwitch s1: lo:127.0.0.1,s1-eth1:None,s1-eth2:None,s1-eth3:None pid=13383>
mininet> net
h1 h1-eth0:s1-eth1
h2 h2-eth0:s1-eth2
h3 h3-eth0:s1-eth3
s1 lo: s1-eth1:h1-eth0 s1-eth2:h2-eth0 s1-eth3:h3-eth0
mininet> sh ovs-ofctl add-flow s1 action=normal
mininet> pingall
*** Ping: testing ping reachability
h1 -> h2 h3
h2 -> h1 h3
h3 -> h1 h2
*** Results: 0% dropped (6/6 received)
mininet>
```

```
File Machine View Input Devices Help

mininet> sh ovs-ofctl show s1
OFPT_FEATURES_REPLY (xid=0x2): dpid:0000000000000001
n_tables:254, n_buffers:0
capabilities: FLOW_STATS TABLE_STATS PORT_STATS QUEUE_STATS ARP_MATCH_IP
actions: output enqueue set_vlan_vid set_vlan_pcp strip_vlan mod_dl_src mod_dl_dst mod_nw_src mod_nw_dst mod_nw_tos mod_tp_src mod_tp_dst
1(s1-eth1): addr:06:15:e2:cb:60:59
  config: 0
  state: 0
  current: 10GB-FD COPPER
  speed: 10000 Mbps now, 0 Mbps max
2(s1-eth2): addr:1a:4f:ee:d6:c9:92
  config: 0
  state: 0
  current: 10GB-FD COPPER
  speed: 10000 Mbps now, 0 Mbps max
3(s1-eth3): addr:82:71:27:a3:54:2e
  config: 0
  state: 0
  current: 10GB-FD COPPER
  speed: 10000 Mbps now, 0 Mbps max
LOCAL(s1): addr:0e:0f:45:de:da:45
  config: PORT_DOWN
  state: LINK_DOWN
  speed: 0 Mbps now, 0 Mbps max
OFPT_GET_CONFIG_REPLY (xid=0x4): frags=normal miss_send_len=0
mininet> sh ovs-ofctl dump-flows s1
 cookie=0x0, duration=214.167s, table=0, n_packets=30, n_bytes=2100, actions=NORMAL
mininet> sh ovs-ofctl del-flows s1
mininet> sh ovs-ofctl dump-flows s1
mininet> pingall
*** Ping: testing ping reachability
h1 -> X X
h2 -> X X
h3 -> X X
*** Results: 100% dropped (0/6 received)
mininet>
```

L1:

File Machine View Input Devices Help

```
h2 -> X X
h3 -> X X
*** Results: 100% dropped (0/6 received)
mininet> sh ovs-ofctl add-flow s1 priority=500,in_port=1,action=output:2
mininet> sh ovs-ofctl adds-flow s1 priority=500,in_port=2,actions=output:1
ovs-ofctl: unknown command 'adds-flow'; use --help for help
mininet> sh ovs-ofctl add-flow s1 priority=500,in_port=2,actions=output:1
mininet> sh ovs-ofctl dump-flows s1
 cookie=0x0, duration=145.493s, table=0, n_packets=0, n_bytes=0, priority=500,in_port="s1-eth1" actions=output:"s1-eth2"
 cookie=0x0, duration=48.538s, table=0, n_packets=0, n_bytes=0, priority=500,in_port="s1-eth2" actions=output:"s1-eth1"
mininet> h1 ping -c2 h2
PING 10.0.0.2 (10.0.0.2) 56(84) bytes of data.
64 bytes from 10.0.0.2: icmp_seq=1 ttl=64 time=0.386 ms
64 bytes from 10.0.0.2: icmp_seq=2 ttl=64 time=0.068 ms

--- 10.0.0.2 ping statistics ---
2 packets transmitted, 2 received, 0% packet loss, time 1032ms
rtt min/avg/max/mdev = 0.068/0.227/0.386/0.159 ms
mininet>
mininet> h3 ping -c2 h2
PING 10.0.0.2 (10.0.0.2) 56(84) bytes of data.
From 10.0.0.3 icmp_seq=1 Destination Host Unreachable
From 10.0.0.3 icmp_seq=2 Destination Host Unreachable

--- 10.0.0.2 ping statistics ---
2 packets transmitted, 0 received, +2 errors, 100% packet loss, time 1009ms
pipe 2
mininet> sh ovs-ofctl add-flow s1 priority=32768,actions=drop
mininet> h1 ping -c2 h2
PING 10.0.0.2 (10.0.0.2) 56(84) bytes of data.

--- 10.0.0.2 ping statistics ---
2 packets transmitted, 0 received, 100% packet loss, time 1014ms

mininet>
```

File Machine View Input Devices Help

```
mininet> h1 ping -c2 h2
PING 10.0.0.2 (10.0.0.2) 56(84) bytes of data.
64 bytes from 10.0.0.2: icmp_seq=1 ttl=64 time=0.386 ms
64 bytes from 10.0.0.2: icmp_seq=2 ttl=64 time=0.068 ms

--- 10.0.0.2 ping statistics ---
2 packets transmitted, 2 received, 0% packet loss, time 1032ms
rtt min/avg/max/mdev = 0.068/0.227/0.386/0.159 ms
mininet>
mininet> h3 ping -c2 h2
PING 10.0.0.2 (10.0.0.2) 56(84) bytes of data.
From 10.0.0.3 icmp_seq=1 Destination Host Unreachable
From 10.0.0.3 icmp_seq=2 Destination Host Unreachable

--- 10.0.0.2 ping statistics ---
2 packets transmitted, 0 received, +2 errors, 100% packet loss, time 1009ms
pipe 2
mininet> sh ovs-ofctl add-flow s1 priority=32768,actions=drop
mininet> h1 ping -c2 h2
PING 10.0.0.2 (10.0.0.2) 56(84) bytes of data.

--- 10.0.0.2 ping statistics ---
2 packets transmitted, 0 received, 100% packet loss, time 1014ms

mininet> sh ovs-ofctl dump-flows s1
 cookie=0x0, duration=144.237s, table=0, n_packets=5, n_bytes=322, actions=drop
 cookie=0x0, duration=477.345s, table=0, n_packets=5, n_bytes=350, priority=500,in_port="s1-eth1" ac
tions=output:"s1-eth2"
 cookie=0x0, duration=380.390s, table=0, n_packets=4, n_bytes=280, priority=500,in_port="s1-eth2" ac
tions=output:"s1-eth1"
mininet> sh ovs-ofctl del-flows s1 --strict
mininet> sh ovs-ofctl dump-flows s1
 cookie=0x0, duration=630.924s, table=0, n_packets=5, n_bytes=350, priority=500,in_port="s1-eth1" ac
tions=output:"s1-eth2"
 cookie=0x0, duration=533.969s, table=0, n_packets=4, n_bytes=280, priority=500,in_port="s1-eth2" ac
tions=output:"s1-eth1"
mininet>
```

```
File Machine View Input Devices Help

mininet>
mininet> h3 ping -c2 h2
PING 10.0.0.2 (10.0.0.2) 56(84) bytes of data.
From 10.0.0.3 icmp_seq=1 Destination Host Unreachable
From 10.0.0.3 icmp_seq=2 Destination Host Unreachable

--- 10.0.0.2 ping statistics ---
2 packets transmitted, 0 received, +2 errors, 100% packet loss, time 1009ms
pipe 2
mininet> sh ovs-ofctl add-flow s1 priority=32768,actions=drop
mininet> h1 ping -c2 h2
PING 10.0.0.2 (10.0.0.2) 56(84) bytes of data.

--- 10.0.0.2 ping statistics ---
2 packets transmitted, 0 received, 100% packet loss, time 1014ms

mininet> sh ovs-ofctl dump-flows s1
 cookie=0x0, duration=144.237s, table=0, n_packets=5, n_bytes=322, actions=drop
 cookie=0x0, duration=477.345s, table=0, n_packets=5, n_bytes=350, priority=500,in_port="s1-eth1" ac
 tions=output:"s1-eth2"
 cookie=0x0, duration=380.390s, table=0, n_packets=4, n_bytes=280, priority=500,in_port="s1-eth2" ac
 tions=output:"s1-eth1"
mininet> sh ovs-ofctl del-flows s1 --strict
mininet> sh ovs-ofctl dump-flows s1
 cookie=0x0, duration=630.924s, table=0, n_packets=5, n_bytes=350, priority=500,in_port="s1-eth1" ac
 tions=output:"s1-eth2"
 cookie=0x0, duration=533.969s, table=0, n_packets=4, n_bytes=280, priority=500,in_port="s1-eth2" ac
 tions=output:"s1-eth1"
mininet> h1 ping -c2 h2
PING 10.0.0.2 (10.0.0.2) 56(84) bytes of data.
64 bytes from 10.0.0.2: icmp_seq=1 ttl=64 time=0.417 ms
64 bytes from 10.0.0.2: icmp_seq=2 ttl=64 time=0.065 ms

--- 10.0.0.2 ping statistics ---
2 packets transmitted, 2 received, 0% packet loss, time 1002ms
rtt min/avg/max/mdev = 0.065/0.241/0.417/0.176 ms
mininet>
```

L2:

File Machine View Input Devices Help

```
cookie=0x0, duration=630.924s, table=0, n_packets=5, n_bytes=350, priority=500,in_port="s1-eth1" actions=output:"s1-eth2"
cookie=0x0, duration=533.969s, table=0, n_packets=4, n_bytes=280, priority=500,in_port="s1-eth2" actions=output:"s1-eth1"
mininet> h1 ping -c2 h2
PING 10.0.0.2 (10.0.0.2) 56(84) bytes of data.
64 bytes from 10.0.0.2: icmp_seq=1 ttl=64 time=0.417 ms
64 bytes from 10.0.0.2: icmp_seq=2 ttl=64 time=0.065 ms

--- 10.0.0.2 ping statistics ---
2 packets transmitted, 2 received, 0% packet loss, time 1002ms
rtt min/avg/max/mdev = 0.065/0.241/0.417/0.176 ms
mininet>
mininet>
mininet> sh ovs-ofctl del-flows s1
mininet> sh ovs-ofctl dump-flows s1
mininet>
mininet> h1 ifconfig -a
h1-eth0: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 10.0.0.1 netmask 255.0.0.0 broadcast 10.255.255.255
    inet6 fe80::200:ff:fe00:1 prefixlen 64 scopeid 0x20<link>
    ether 00:00:00:00:00:01 txqueuelen 1000 (Ethernet)
    RX packets 37 bytes 2678 (2.6 KB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 44 bytes 3096 (3.0 KB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536
    inet 127.0.0.1 netmask 255.0.0.0
    inet6 ::1 prefixlen 128 scopeid 0x10<host>
    loop txqueuelen 1000 (Local Loopback)
    RX packets 0 bytes 0 (0.0 B)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 0 bytes 0 (0.0 B)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

mininet> _
```




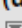
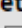

```
    inet 10.0.0.2 netmask 255.0.0.0 broadcast 10.255.255.255
    inet6 fe80::200:ff:fe00:2 prefixlen 64 scopeid 0x20<link>
    ether 00:00:00:00:00:02 txqueuelen 1000 (Ethernet)
    RX packets 37 bytes 2678 (2.6 KB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 39 bytes 2774 (2.7 KB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536
    inet 127.0.0.1 netmask 255.0.0.0
    inet6 ::1 prefixlen 128 scopeid 0x10<host>
    loop txqueuelen 1000 (Local Loopback)
    RX packets 0 bytes 0 (0.0 B)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 0 bytes 0 (0.0 B)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

mininet> h3 ifconfig -a
h3-eth0: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 10.0.0.3 netmask 255.0.0.0 broadcast 10.255.255.255
    inet6 fe80::200:ff:fe00:3 prefixlen 64 scopeid 0x20<link>
    ether 00:00:00:00:00:03 txqueuelen 1000 (Ethernet)
    RX packets 29 bytes 2118 (2.1 KB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 34 bytes 2340 (2.3 KB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536
    inet 127.0.0.1 netmask 255.0.0.0
    inet6 ::1 prefixlen 128 scopeid 0x10<host>
    loop txqueuelen 1000 (Local Loopback)
    RX packets 2 bytes 224 (224.0 B)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 2 bytes 224 (224.0 B)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

mininet> _
```

| Ethertype (decimal)  | Ethertype (hex)  | Exp. Ethernet (decimal)  | Exp. Ethernet (octal)  | Description  |
|----------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------|
| 0000 | 0000-05DC | - | - | IEEE802.3 Length Field |
| 0257 | 0101-01FF | - | - | Experimental |
| 0512 | 0200 | 512 | 1000 | XEROX PUP (see 0A00) |
| 0513 | 0201 | - | - | PUP Addr Trans (see 0A01) |
| | 0400 | | | Nixdorf |
| 1536 | 0600 | 1536 | 3000 | XEROX NS IDP |
| | | | | |
| | 0660 | | | DLOG |
| | 0661 | | | DLOG |
| 2048 | 0800 | 513 | 1001 | Internet Protocol version 4 (IPv4) |
| 2049 | 0801 | - | - | X.75 Internet |
| 2050 | 0802 | - | - | NBS Internet |
| 2051 | 0803 | - | - | ECMA Internet |
| 2052 | 0804 | - | - | Chaosnet |
| 2053 | 0805 | - | - | X.25 Level 3 |
| 2054 | 0806 | - | - | Address Resolution Protocol (ARP) |
| 2055 | 0807 | - | - | XNS Compatability |
| 2056 | 0808 | - | - | Frame Relay ARP |
| 2076 | 081C | - | - | Symbolics Private |
| 2184 | 0888-088A | - | - | Xyplex |
| 2204 | 0800 | | | |

```
mininet> h3 ifconfig -a
h3-eth0: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 10.0.0.3 netmask 255.0.0.0 broadcast 10.255.255.255
    inet6 fe80::200:ff:fe00:3 prefixlen 64 scopeid 0x20<link>
    ether 00:00:00:00:00:03 txqueuelen 1000 (Ethernet)
    RX packets 29 bytes 2118 (2.1 KB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 34 bytes 2340 (2.3 KB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536
    inet 127.0.0.1 netmask 255.0.0.0
    inet6 ::1 prefixlen 128 scopeid 0x10<host>
    loop txqueuelen 1000 (Local Loopback)
    RX packets 2 bytes 224 (224.0 B)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 2 bytes 224 (224.0 B)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

mininet>
mininet> sh ovs-ofctl add-flow s1 dl_src=00:00:00:00:00:01,dl_dst=00:00:00:00:00:02,actions=output:2
mininet> sh ovs-ofctl add-flow s1 dl_src=00:00:00:00:00:02,dl_dst=00:00:00:00:00:01,actions=output:1
mininet> sh ovs-ofctl add-flow s1 dl_type=0x806,nw_proto=1,actions=flood
mininet> sh ovs-ofctl dump-flows s1
    cookie=0x0, duration=235.517s, table=0, n_packets=0, n_bytes=0, dl_src=00:00:00:00:00:01,dl_dst=00:
00:00:00:00:00:02 actions=output:"s1-eth2"
    cookie=0x0, duration=147.870s, table=0, n_packets=0, n_bytes=0, dl_src=00:00:00:00:00:02,dl_dst=00:
00:00:00:00:00:01 actions=output:"s1-eth1"
    cookie=0x0, duration=25.402s, table=0, n_packets=0, n_bytes=0, arp,arp_op=1 actions=FL00D
mininet> pingall
*** Ping: testing ping reachability
h1 -> h2 X
h2 -> h1 X
h3 -> X X
*** Results: 66% dropped (2/6 received)
mininet> _
```

L3:

```
mininet> pingall
*** Ping: testing ping reachability
h1 -> h2 X
h2 -> h1 X
h3 -> X X
*** Results: 66% dropped (2/6 received)
mininet>
mininet> sh ovs-ofctl del-flows s1
mininet> sh ovs-ofctl dump-flows s1
mininet>
mininet> sh ovs-ofctl add-flow s1 priority=500,dl_type=0x800,nw_src=10.0.0.0/24,nw_dst=10.0.0.0/24,actions=normal
mininet> sh ovs-ofctl add-flow s1 priority=800,ip,nw_src=10.0.0.3,actions=mod_nw_tos:184,normal
mininet> sh ovs-ofctl add-flow s1 arp,nw_dst=10.0.0.1,actions=output:1
mininet> sh ovs-ofctl add-flow s1 arp,nw_dst=10.0.0.2,actions=output:2
mininet> sh ovs-ofctl add-flow s1 arp,nw_dst=10.0.0.3,actions=output:3
mininet> pingall
*** Ping: testing ping reachability
h1 -> h2 h3
h2 -> h1 h3
h3 -> h1 h2
*** Results: 0% dropped (6/6 received)
mininet> sh ovs-ofctl dump-flows s1
    cookie=0x0, duration=180.199s, table=0, n_packets=4, n_bytes=168, arp,arp_tpa=10.0.0.1 actions=output:"s1-eth1"
    cookie=0x0, duration=122.974s, table=0, n_packets=4, n_bytes=168, arp,arp_tpa=10.0.0.2 actions=output:"s1-eth2"
    cookie=0x0, duration=68.424s, table=0, n_packets=4, n_bytes=168, arp,arp_tpa=10.0.0.3 actions=output:"s1-eth3"
    cookie=0x0, duration=287.718s, table=0, n_packets=4, n_bytes=392, priority=800,ip,nw_src=10.0.0.3 actions=mod_nw_tos:184,NORMAL
    cookie=0x0, duration=520.910s, table=0, n_packets=8, n_bytes=784, priority=500,ip,nw_src=10.0.0.0/24,nw_dst=10.0.0.0/24 actions=NORMAL
mininet> sh ovs-ofctl del-flows s1
mininet>
mininet> sh ovs-ofctl dump-flows s1
mininet> _
```

L4:

```
cookie=0x0, duration=520.910s, table=0, n_packets=8, n_bytes=784, priority=500,ip,nw_src=10.0.0.0/24,nw_dst=10.0.0.0/24 actions=NORMAL
mininet> sh ovs-ofctl del-flows s1
mininet>
mininet> sh ovs-ofctl dump-flows s1
mininet> h3 python -m SimpleHTTPServer 80 &
mininet> sh ovs-ofctl add-flow s1 actions=normal
mininet> h1 curl h3
<!DOCTYPE html PUBLIC "-//W3C//DTD HTML 3.2 Final//EN"><html>
<title>Directory listing for /</title>
<body>
<h2>Directory listing for /</h2>
<hr>
<ul>
<li><a href=".bash_history">.bash_history</a>
<li><a href=".bash_logout">.bash_logout</a>
<li><a href=".bashrc">.bashrc</a>
<li><a href=".cache/">.cache/</a>
<li><a href=".gnupg/">.gnupg/</a>
<li><a href=".karaf/">.karaf/</a>
<li><a href=".local/">.local/</a>
<li><a href=".mininet_history">.mininet_history</a>
<li><a href=".profile">.profile</a>
<li><a href=".ssh/">.ssh/</a>
<li><a href=".sudo_as_admin_successful">.sudo_as_admin_successful</a>
<li><a href=".wget-hsts">.wget-hsts</a>
<li><a href="distribution-karaf-0.6.4-Carbon/">distribution-karaf-0.6.4-Carbon/</a>
<li><a href="distribution-karaf-0.6.4-Carbon.zip">distribution-karaf-0.6.4-Carbon.zip</a>
<li><a href="mininet/">mininet/</a>
<li><a href="ryu/">ryu/</a>
<li><a href="v-ryu/">v-ryu/</a>
<li><a href="vxlan-lab/">vxlan-lab/</a>
</ul>
<hr>
</body>
</html>
mininet> _
```

```
File Machine View Input Devices Help

<li><a href="ryu/">ryu/</a>
<li><a href="v-ryu/">v-ryu/</a>
<li><a href="vxlan-lab/">vxlan-lab/</a>
</ul>
<hr>
</body>
</html>
mininet> h2 curl h3
<!DOCTYPE html PUBLIC "-//W3C//DTD HTML 3.2 Final//EN"><html>
<title>Directory listing for /</title>
<body>
<h2>Directory listing for /</h2>
<hr>
<ul>
<li><a href=".bash_history">.bash_history</a>
<li><a href=".bash_logout">.bash_logout</a>
<li><a href=".bashrc">.bashrc</a>
<li><a href=".cache/">.cache/</a>
<li><a href=".gnupg/">.gnupg/</a>
<li><a href=".karaf/">.karaf/</a>
<li><a href=".local/">.local/</a>
<li><a href=".mininet_history">.mininet_history</a>
<li><a href=".profile">.profile</a>
<li><a href=".ssh/">.ssh/</a>
<li><a href=".sudo_as_admin_successful">.sudo_as_admin_successful</a>
<li><a href=".wget-hsts">.wget-hsts</a>
<li><a href="distribution-karaf-0.6.4-Carbon/">distribution-karaf-0.6.4-Carbon/</a>
<li><a href="distribution-karaf-0.6.4-Carbon.zip">distribution-karaf-0.6.4-Carbon.zip</a>
<li><a href="mininet/">mininet/</a>
<li><a href="ryu/">ryu/</a>
<li><a href="v-ryu/">v-ryu/</a>
<li><a href="vxlan-lab/">vxlan-lab/</a>
</ul>
<hr>
</body>
</html>
mininet>
```

Open vSwitch Multiple Flow Tables

```
File Machine View Input Devices Help
ks-sdn@sdn_server:~$ sudo mn --topo=single,3 --controller=none --mac
[sudo] password for ks-sdn:
*** Creating network
*** Adding controller
*** Adding hosts:
h1 h2 h3
*** Adding switches:
s1
*** Adding links:
(h1, s1) (h2, s1) (h3, s1)
*** Configuring hosts
h1 h2 h3
*** Starting controller

*** Starting 1 switches
s1 ...
*** Starting CLI:
mininet> h1 route add default gw 10.0.0.254 h1-eth0
mininet> h1 arp -s 10.0.0.254 00:00:00:00:11:11
mininet> h2 arp -s 10.0.0.254 00:00:00:00:11:11
mininet> h3 ifconfig h3-eth0 30.0.0.3 netmask 255.255.255.0
mininet> h3 route add default gw 30.0.0.254 h3-eth0
mininet> h3 arp -s 30.0.0.254 00:00:00:00:33:33
mininet> h3 sudo python -m SimpleHTTPServer 80 &
mininet>
```

```
File Machine View Input Devices Help
GNU nano 2.9.3 tables.txt Modified

#table 0 - Access Control
table=0,ip,nw_src=10.0.0.0/24,nw_dst=10.0.0.0/24,actions=resubmit(,1)
table=0,arp,nw_src=10.0.0.0/24,nw_dst=10.0.0.0/24,actions=resubmit(,1)
table=0,icmp,nw_src=10.0.0.1,nw_dst=30.0.0.3,actions=resubmit(,1)
table=0,tcp,nw_src=10.0.0.1,nw_dst=30.0.0.3,tp_dst=80,actions=resubmit(,1)
table=0,ip,nw_src=30.0.0.3,actions=resubmit(,1)
table=0,priority=0,actions=drop

#table 1 - NAT
table=1,ip,nw_src=10.0.0.1,nw_dst=30.0.0.3,actions=mod_nw_src=5.5.5.5,resubmit(,2)
table=1,ip,nw_src=30.0.0.3,nw_dst=5.5.5.5,actions=mod_nw_dst=10.0.0.1,resubmit(,2)
table=1,priority=0,actions=resubmit(,2)

#table 2 forward/route
table=2,ip,nw_dst=10.0.0.1,actions=mod_dl_dst=00:00:00:00:00:01,output:1
table=2,ip,nw_dst=10.0.0.2,actions=mod_dl_dst=00:00:00:00:00:02,output:2
table=2,ip,nw_dst=30.0.0.3,actions=mod_dl_dst=00:00:00:00:00:03,output:3
priority=0,table=2,arp,nw_dst=10.0.0.1,actions=output:1
priority=0,table=2,arp,nw_dst=10.0.0.2,actions=output:2
-

^G Get Help ^O Write Out ^W Where Is ^K Cut Text ^J Justify ^C Cur Pos M-U Undo
^X Exit ^R Read File ^_ Replace ^U Uncut Text ^T To Spell ^_ Go To Line M-E Redo
```

```
mininet>
mininet>
mininet> sh ovs-ofctl dump-flows s1 -O OpenFlow13 > exampleLog.txt
mininet> sh more tables.txt
#table 0 - Access Control
table=0,ip,nw_src=10.0.0.0/24,nw_dst=10.0.0.0/24,actions=resubmit(,1)
table=0,arp,nw_src=10.0.0.0/24,nw_dst=10.0.0.0/24,actions=resubmit(,1)
table=0,icmp,nw_src=10.0.0.1,nw_dst=30.0.0.3,actions=resubmit(,1)
table=0,tcp,nw_src=10.0.0.1,nw_dst=30.0.0.3,tp_dst=80,actions=resubmit(,1)
table=0,ip,nw_src=30.0.0.3,actions=resubmit(,1)
table=0,priority=0,actions=drop

#table 1 - NAT
table=1,ip,nw_src=10.0.0.1,nw_dst=30.0.0.3,actions=mod_nw_src=5.5.5.5,resubmit(,2)
table=1,ip,nw_src=30.0.0.3,nw_dst=5.5.5.5,actions=mod_nw_dst=10.0.0.1,resubmit(,2)
table=1,priority=0,actions=resubmit(,2)

#table 2 forward/route
table=2,ip,nw_dst=10.0.0.1,actions=mod_dl_dst=00:00:00:00:00:01,output:1
table=2,ip,nw_dst=10.0.0.2,actions=mod_dl_dst=00:00:00:00:00:02,output:2
table=2,ip,nw_dst=30.0.0.3,actions=mod_dl_dst=00:00:00:00:00:03,output:3
priority=0,table=2,arp,nw_dst=10.0.0.1,actions=output:1
priority=0,table=2,arp,nw_dst=10.0.0.2,actions=output:2

mininet>
```



```
#table 2 forward/route
table=2,ip,nw_dst=10.0.0.1,actions=mod_dl_dst=00:00:00:00:00:01,output:1
table=2,ip,nw_dst=10.0.0.2,actions=mod_dl_dst=00:00:00:00:00:02,output:2
table=2,ip,nw_dst=30.0.0.3,actions=mod_dl_dst=00:00:00:00:00:03,output:3
priority=0,table=2,arp,nw_dst=10.0.0.1,actions=output:1
priority=0,table=2,arp,nw_dst=10.0.0.2,actions=output:2
```

```
mininet> sh ovs-ofctl add-flows s1 tables.txt
mininet> h1 ping -c2 h2
PING 10.0.0.2 (10.0.0.2) 56(84) bytes of data.
64 bytes from 10.0.0.2: icmp_seq=1 ttl=64 time=0.495 ms
64 bytes from 10.0.0.2: icmp_seq=2 ttl=64 time=0.069 ms

--- 10.0.0.2 ping statistics ---
2 packets transmitted, 2 received, 0% packet loss, time 1031ms
rtt min/avg/max/mdev = 0.069/0.282/0.495/0.213 ms
mininet> h1 ping -c2 30.0.0.3
PING 30.0.0.3 (30.0.0.3) 56(84) bytes of data.
64 bytes from 30.0.0.3: icmp_seq=1 ttl=64 time=0.365 ms
64 bytes from 30.0.0.3: icmp_seq=2 ttl=64 time=0.056 ms

--- 30.0.0.3 ping statistics ---
2 packets transmitted, 2 received, 0% packet loss, time 1001ms
rtt min/avg/max/mdev = 0.056/0.210/0.365/0.155 ms
mininet> h2 ping -c2 30.0.0.3
connect: Network is unreachable
mininet>
```

```
File Machine View Input Devices Help
--- 30.0.0.3 ping statistics ---
2 packets transmitted, 2 received, 0% packet loss, time 1001ms
rtt min/avg/max/mdev = 0.056/0.210/0.365/0.155 ms
mininet> h2 ping -c2 30.0.0.3
connect: Network is unreachable
mininet> h1 curl 30.0.0.3
<!DOCTYPE html PUBLIC "-//W3C//DTD HTML 3.2 Final//EN"><html>
<title>Directory listing for /</title>
<body>
<h2>Directory listing for /</h2>
<hr>
<ul>
<li><a href=".bash_history">.bash_history</a>
<li><a href=".bash_logout">.bash_logout</a>
<li><a href=".bashrc">.bashrc</a>
<li><a href=".cache/">.cache</a>
<li><a href=".gnupg/">.gnupg</a>
<li><a href=".karaf/">.karaf</a>
<li><a href=".local/">.local</a>
<li><a href=".mininet_history">.mininet_history</a>
<li><a href=".profile">.profile</a>
<li><a href=".ssh/">.ssh</a>
<li><a href=".sudo_as_admin_successful">.sudo_as_admin_successful</a>
<li><a href=".wget-hsts">.wget-hsts</a>
<li><a href="distribution-karaf-0.6.4-Carbon/">distribution-karaf-0.6.4-Carbon</a>
<li><a href="distribution-karaf-0.6.4-Carbon.zip">distribution-karaf-0.6.4-Carbon.zip</a>
<li><a href="exampleLog.txt">exampleLog.txt</a>
<li><a href="mininet/">mininet</a>
<li><a href="ryu/">ryu</a>
<li><a href="tables.txt">tables.txt</a>
<li><a href="v-ryu/">v-ryu</a>
<li><a href="vxlan-lab/">vxlan-lab</a>
</ul>
<hr>
</body>
</html>
mininet>
```

Hands-on Training with Python Net Apps - Elective 2

APP - 1

```
ks-sdn@sdn_server:~/mininet/examples$ ls
baresshd.py      controllers2.py  limit.py        multitest.py    scratchnetuser.py
bind.py          controllers.py   linearbandwidth.py natnet.py        simpleperf.py
clustercli.py     controlnet.py   linuxrouter.py  nat.py          sshd.py
clusterdemo.py    cpu.py         miniedit.py     numberedports.py test
clusterperf.py    emptynet.py    mobility.py     popenpoll.py    tree1024.py
cluster.py        hwintf.py      multilink.py    popen.py        treeping64.py
clusterSanity.py  __init__.py     multiping.py    README.md       vlanhost.py
consoles.py       intfoptions.py  multipoll.py    scratchnet.py
ks-sdn@sdn_server:~/mininet/examples$ cat controllers.py _
```

```
"""
Create a network where different switches are connected to
different controllers, by creating a custom Switch() subclass.
"""

from mininet.net import Mininet
from mininet.node import OVSSwitch, Controller, RemoteController
from mininet.topolib import TreeTopo
from mininet.log import setLogLevel
from mininet.cli import CLI

setLogLevel( 'info' )

# Two local and one "external" controller (which is actually c0)
# Ignore the warning message that the remote isn't (yet) running
c0 = Controller( 'c0', port=6633 )
c1 = Controller( 'c1', port=6634 )
c2 = RemoteController( 'c2', ip='127.0.0.1', port=6633 )

cmap = { 's1': c0, 's2': c1, 's3': c2 }

class MultiSwitch( OVSSwitch ):
    "Custom Switch() subclass that connects to different controllers"
    def start( self, controllers ):
        return OVSSwitch.start( self, [ cmap[ self.name ] ] )

topo = TreeTopo( depth=2, fanout=2 )
net = Mininet( topo=topo, switch=MultiSwitch, build=False, waitConnected=True )
for c in [ c0, c1 ]:
    net.addController(c)
net.build()
net.start()
CLI( net )
net.stop()
ks-sdn@sdn_server:~/mininet/examples$ _
```

Explanation of Code App:

- The first statement is a string literal, this is the module documentation string, explaining what the app does, with the help of docstring tools, user can interactively browse through the code.
- The five import statements are used to import **Mininet**, **OVSSwitch**, **Controller**, **RemoteController**, **TreeTopo**, **setLogLevel**, and **CLI** modules from their parent modules respectively.
- **setLogLevel()** function was called by passing an argument to setup log level.
- **#** is used to write comments for the code
- Two instances of class **Controller**, **c0**, **c1** were instantiated with respective instance variables, to run executable files (OpenFlow Controllers)
- one instance of class **RemoteController** were also instantiated with its instance variables.
- A dictionary of the instances was assigned to **cmap** in the form of key:value pairs enclosed with **{}** to form a data structure where the unique keys are **s1**, **s2**, and **s3** and values are **c0**, **c1**, **c2**.
- A **MultiSwitch** class was created with a valid attribute (function) to start up an OVS OpenFlow Switches on controllers.
- Instance of **TreeTopo** class with instance variables (**depth=2**, **fanout=2**) instantiated
- Instance of **Mininet** class with instance variables instantiated
- The **for** statement is used to iterate a list and then add controller, this will add two controllers.
- **net.build()** build mininet network.
- **net.start()** start controllers and switches.
- **CLI(net)** runs a batch or interactive mode for the network.
- The **net.stop()** stop controllers, switches, and hosts.

RESULT OF CODE:

```

ks-sdn@sdn_server:~/mininet/examples$ ls
baresshd.py      controllers2.py  limit.py        multitest.py    scratchnetuser.py
bind.py          controllers.py  linearbandwidth.py  natnet.py       simpleperf.py
clustercli.py    controlnet.py  linuxrouter.py   nat.py          sshd.py
clusterdemo.py   cpu.py         miniedit.py      numberedports.py  test
clusterperf.py   emptynet.py    mobility.py      popenpoll.py    tree1024.py
cluster.py        hwintf.py      multilink.py     popen.py        treeping64.py
clusterSanity.py  __init__.py    multiping.py     README.md       vlanhost.py
consoles.py      intfoptions.py  multipoll.py     scratchnet.py

ks-sdn@sdn_server:~/mininet/examples$ cd ../../
ks-sdn@sdn_server:~$ sudo mn -c
[sudo] password for ks-sdn:
*** Removing excess controllers/ofprotocols/ofdatapaths/pings/noxes
killall controller ofprotocol ofdatapath ping nox_corelt-nox_core ovs-openflowd ovs-controllerovs-te
stcontroller udpbwtest mnexec ivs ryu-manager 2> /dev/null
killall -9 controller ofprotocol ofdatapath ping nox_corelt-nox_core ovs-openflowd ovs-controllerovs
-testcontroller udpbwtest mnexec ivs ryu-manager 2> /dev/null
pkill -9 -f "sudo mnexec"
*** Removing junk from /tmp
rm -f /tmp/vconn* /tmp/vlogs* /tmp/*.out /tmp/*.log
*** Removing old X11 tunnels
*** Removing excess kernel datapaths
ps ax | egrep -o 'dp[0-9]+' | sed 's/dp/nl:/'
*** Removing OVS datapaths
ovs-vsctl --timeout=1 list-br
ovs-vsctl --timeout=1 list-br
*** Removing all links of the pattern foo-ethX
ip link show | egrep -o '([_-[:alnum:]]+-eth[[:digit:]]+)'
ip link show
*** Killing stale mininet node processes
pkill -9 -f mininet:
*** Shutting down stale tunnels
pkill -9 -f Tunnel=Ethernet
pkill -9 -f .ssh/mn
rm -f ~/.ssh/mn/*
*** Cleanup complete.
ks-sdn@sdn_server:~$

```

```

ks-sdn@sdn_server:~$ sudo python ./mininet/examples/controllers.py
Unable to contact the remote controller at 127.0.0.1:6633
*** Creating network
*** Adding hosts:
h1 h2 h3 h4
*** Adding switches:
s1 s2 s3
*** Adding links:
(s1, s2) (s1, s3) (s2, h1) (s2, h2) (s3, h3) (s3, h4)
*** Configuring hosts
h1 h2 h3 h4
*** Starting controller
c0 c1
*** Starting 3 switches
s1 s2 s3 ...
*** Waiting for switches to connect
s1 s2 s3
*** Starting CLI:
mininet> _

```

```

*** Waiting for switches to connect
s1 s2 s3
*** Starting CLI:
mininet> help

Documented commands (type help <topic>):
=====
EOF      gterm  iperfudp  nodes      pingpair    py      switch  xterm
dpctl    help   link     noecho     pingpairfull  quit    time
dump     intf  links    pingall    ports       sh      wait
exit     iperf  net      pingallfull  px          source  x

You may also send a command to a node using:
  <node> command {args}
For example:
  mininet> h1 ifconfig

The interpreter automatically substitutes IP addresses
for node names when a node is the first arg, so commands
like
  mininet> h2 ping h3
should work.

Some character-oriented interactive commands require
noecho:
  mininet> noecho h2 vi foo.py
However, starting up an xterm/gterm is generally better:
  mininet> xterm h2

mininet> pingall
*** Ping: testing ping reachability
h1 -> h2 h3 h4
h2 -> h1 h3 h4
h3 -> h1 h2 h4
h4 -> h1 h2 h3
*** Results: 0% dropped (12/12 received)
mininet> _

```

```

h3 -> h1 h2 h4
h4 -> h1 h2 h3
*** Results: 0% dropped (12/12 received)
mininet> dump
<Host h1: h1-eth0:10.0.0.1 pid=2920>
<Host h2: h2-eth0:10.0.0.2 pid=2922>
<Host h3: h3-eth0:10.0.0.3 pid=2924>
<Host h4: h4-eth0:10.0.0.4 pid=2926>
<MultiSwitch s1: lo:127.0.0.1,s1-eth1:None,s1-eth2:None pid=2931>
<MultiSwitch s2: lo:127.0.0.1,s2-eth1:None,s2-eth2:None,s2-eth3:None pid=2934>
<MultiSwitch s3: lo:127.0.0.1,s3-eth1:None,s3-eth2:None,s3-eth3:None pid=2937>
<Controller c0: 127.0.0.1:6633 pid=2903>
<Controller c1: 127.0.0.1:6634 pid=2908>
mininet> nodes
available nodes are:
c0 c1 h1 h2 h3 h4 s1 s2 s3
mininet> net
h1 h1-eth0:s2-eth1
h2 h2-eth0:s2-eth2
h3 h3-eth0:s3-eth1
h4 h4-eth0:s3-eth2
s1 lo: s1-eth1:s2-eth3 s1-eth2:s3-eth3
s2 lo: s2-eth1:h1-eth0 s2-eth2:h2-eth0 s2-eth3:s1-eth1
s3 lo: s3-eth1:h3-eth0 s3-eth2:h4-eth0 s3-eth3:s1-eth2
c0
c1
mininet> exit
*** Stopping 2 controllers
c0 c1
*** Stopping 6 links
.....
*** Stopping 3 switches
s1 s2 s3
*** Stopping 4 hosts
h1 h2 h3 h4
*** Done
ks-sdn@sdn_server:~$

```

APP - 2

```
#!/usr/bin/env python
```

```
"""
```

```
This example creates a multi-controller network from semi-scratch by  
using the net.add*() API and manually starting the switches and controllers.
```

```
This is the "mid-level" API, which is an alternative to the "high-level"  
Topo() API which supports parametrized topology classes.
```

```
Note that one could also create a custom switch class and pass it into  
the Mininet() constructor.
```

```
"""
```

```
from mininet.net import Mininet  
from mininet.node import Controller, OVSSwitch  
from mininet.cli import CLI  
from mininet.log import setLogLevel, info
```

```
def multiControllerNet():  
    "Create a network from semi-scratch with multiple controllers."
```

```
    net = Mininet( controller=Controller, switch=OVSSwitch,  
                   waitConnected=True )
```

```
    info( "*** Creating (reference) controllers\n" )
```

```
    c1 = net.addController( 'c1', port=6633 )
```

```
    c2 = net.addController( 'c2', port=6634 )
```

```
    info( "*** Creating switches\n" )
```

```
    s1 = net.addSwitch( 's1' )
```

```
    s2 = net.addSwitch( 's2' )
```

```
^G Get Help  
^X Exit
```

```
^O Write Out  
^R Read File
```

```
^W Where Is  
^_ Replace
```

```
^K Cut Text  
^U Uncut Text
```

```
^J Justify  
^T To Linter
```

```
^C Cur Pos  
^_ Go To Line
```

```
M-U Undo  
M-E Redo
```

```
GNU nano 2.9.3                               ./mininet/examples/controllers2.py

s2 = net.addSwitch( 's2' )

info( "*** Creating hosts\n" )
hosts1 = [ net.addHost( 'h%d' % n ) for n in ( 3, 4 ) ]
hosts2 = [ net.addHost( 'h%d' % n ) for n in ( 5, 6 ) ]

info( "*** Creating links\n" )
for h in hosts1:
    net.addLink( s1, h )
for h in hosts2:
    net.addLink( s2, h )
net.addLink( s1, s2 )

info( "*** Starting network\n" )
net.build()
c1.start()
c2.start()
s1.start( [ c1 ] )
s2.start( [ c2 ] )

info( "*** Testing network\n" )
net.pingAll()

info( "*** Running CLI\n" )
CLI( net )

info( "*** Stopping network\n" )
net.stop()

if __name__ == '__main__':
    setLogLevel( 'info' ) # for CLI output

^G Get Help  ^O Write Out  ^W Where Is   ^K Cut Text   ^J Justify    ^C Cur Pos    M-U Undo
^X Exit      ^R Read File  ^_ Replace    ^U Uncut Text ^I To Linter  ^_ Go To Line M-E Redo

info( "*** Stopping network\n" )
net.stop()

if __name__ == '__main__':
    setLogLevel( 'info' ) # for CLI output
    multiControllerNet()

^G Get Help  ^O Write Out  ^W Where Is   ^K Cut Text   ^J Justify    ^C Cur Pos    M-U Undo
^X Exit      ^R Read File  ^_ Replace    ^U Uncut Text ^I To Linter  ^_ Go To Line M-E Redo
```

Code Explanation:

- The first statement is a string literal, this is the module documentation string, explaining what the app does, with the help of docstring tools, user can interactively browse through the code.
- The four import statements are used to import **Mininet**, **OVSSwitch**, **Controller**, **setLogLevel**, **Info**, and **CLI** modules from their parent modules respectively.
- **def** is used to define **multiControllerNet** function that create multiple controllers by first instantiating Mininet class, called **info()** method, added two controllers to local variables **c1**, **c2**, added two switches to local variables **s1**, **s2**, added two hosts to local variables **host1**, and **host2**, used the **for** loop to add links to hosts and switches.
 - **net.build()** build mininet network
 - **c1.start()** start first controller
 - **c2.start()** start second controller
 - **s1.start([c1])** start first switch on first controller
 - **s2.start([c2])** start second switch on second controller
 - **net.pingAll()** pings all hosts
 - **CLI(net)** runs a batch or interactive mode for the network.
 - The **net.stop()** stop controllers, switches, and hosts.
- The **if __name__ == "__main__":** statement makes the file usable as a script and importable module.

This app should create mininet network with two controllers, two switches, and two hosts on each switch, and then ping all nodes.

RESULT:

```
File Machine View Input Devices Help
ks-sdn@sdn_server:~$ sudo python ./mininet/examples/controllers2.py
[sudo] password for ks-sdn:
*** Creating (reference) controllers
*** Creating switches
*** Creating hosts
*** Creating links
*** Starting network
*** Configuring hosts
h3 h4 h5 h6
*** Testing network
*** Ping: testing ping reachability
h3 -> h4 h5 h6
h4 -> h3 h5 h6
h5 -> h3 h4 h6
h6 -> h3 h4 h5
*** Results: 0% dropped (12/12 received)
*** Running CLI
*** Starting CLI:
mininet> dump
<Host h3: h3-eth0:10.0.0.1 pid=3375>
<Host h4: h4-eth0:10.0.0.2 pid=3377>
<Host h5: h5-eth0:10.0.0.3 pid=3379>
<Host h6: h6-eth0:10.0.0.4 pid=3381>
<OVSSwitch s1: lo:127.0.0.1,s1-eth1:None,s1-eth2:None,s1-eth3:None pid=3367>
<OVSSwitch s2: lo:127.0.0.1,s2-eth1:None,s2-eth2:None,s2-eth3:None pid=3370>
<Controller c1: 127.0.0.1:6633 pid=3354>
<Controller c2: 127.0.0.1:6634 pid=3359>
mininet> net
h3 h3-eth0:s1-eth1
h4 h4-eth0:s1-eth2
h5 h5-eth0:s2-eth1
h6 h6-eth0:s2-eth2
s1 lo: s1-eth1:h3-eth0 s1-eth2:h4-eth0 s1-eth3:s2-eth3
s2 lo: s2-eth1:h5-eth0 s2-eth2:h6-eth0 s2-eth3:s1-eth3
c1
c2
mininet> _
```

```

File Machine View Input Devices Help
mininet> exit
*** Stopping network
*** Stopping 2 controllers
c1 c2
*** Stopping 5 links
.....
*** Stopping 2 switches
s1 s2
*** Stopping 4 hosts
h3 h4 h5 h6
*** Done
ks-sdn@sdn_server:~$ sudo mn -c
*** Removing excess controllers/ofprotocols/ofdatapaths/pings/noxes
killall controller ofprotocol ofdatapath ping nox_corelt-nox_core ovs-openflowd ovs-controllerovs-te
stcontroller udpbwtest mnexec ivs ryu-manager 2> /dev/null
killall -9 controller ofprotocol ofdatapath ping nox_corelt-nox_core ovs-openflowd ovs-controllerovs
-testcontroller udpbwtest mnexec ivs ryu-manager 2> /dev/null
pkill -9 -f "sudo mnexec"
*** Removing junk from /tmp
rm -f /tmp/vconn* /tmp/vlogs* /tmp/*.out /tmp/*.log
*** Removing old X11 tunnels
*** Removing excess kernel datapaths
ps ax | egrep -o 'dp[0-9]+' | sed 's/dp/nl:/'
*** Removing OVS datapaths
ovs-vsctl --timeout=1 list-br
ovs-vsctl --timeout=1 list-br
*** Removing all links of the pattern foo-ethX
ip link show | egrep -o '([-_.[:alnum:]]+-eth[[:digit:]]+)'
ip link show
*** Killing stale mininet node processes
pkill -9 -f mininet:
*** Shutting down stale tunnels
pkill -9 -f Tunnel=Ethernet
pkill -9 -f .ssh/mn
rm -f ~/.ssh/mn/*
*** Cleanup complete.
ks-sdn@sdn_server:~$ _

```

APP - 3

```
#!/usr/bin/env python
```

```
"""
```

```
cpu.py: test iperf bandwidth for varying cpu limits
```

Since we are limiting the hosts (only), we should expect the iperf processes to be affected, as well as any system processing which is billed to the hosts.

We reserve >50% of cycles for system processing; we assume that this is enough for it not to affect results. Hosts are limited to 40% of total cycles, which we assume is enough to make them CPU bound.

As CPU performance increases over time, we may have to reduce the overall CPU allocation so that the host processing is still CPU bound. This is perhaps an argument for specifying performance in a more system-independent manner.

It would also be nice to have a better handle on limiting packet processing cycles. It's not entirely clear to me how those are billed to user or system processes if we are using OVS with a kernel datapath. With a user datapath, they are easier to account for, but overall performance is usually lower.

Although the iperf client uses more CPU and should be CPU bound (?), we measure the received data at the server since the client transmit rate includes buffering.

```
"""
```

```
from mininet.net import Mininet
```

```
from mininet.node import CPULimitedHost
```

```
[ Read 107 lines ]
```

```
^G Get Help
```

```
^O Write Out
```

```
^W Where Is
```

```
^K Cut Text
```

```
^J Justify
```

```
^C Cur Pos
```

```
M-U Undo
```

```
^X Exit
```

```
^R Read File
```

```
^_ Replace
```

```
^U Uncut Text
```

```
^T To Linter
```

```
^_ Go To Line
```

```
M-E Redo
```

overall CPU allocation so that the host processing is still CPU bound. This is perhaps an argument for specifying performance in a more system-independent manner.

It would also be nice to have a better handle on limiting packet processing cycles. It's not entirely clear to me how those are billed to user or system processes if we are using OVS with a kernel datapath. With a user datapath, they are easier to account for, but overall performance is usually lower.

Although the iperf client uses more CPU and should be CPU bound (?), we measure the received data at the server since the client transmit rate includes buffering.

```
from mininet.net import Mininet
from mininet.node import CPULimitedHost
from mininet.topolib import TreeTopo
from mininet.util import custom, waitListening, decode
from mininet.log import setLogLevel, info
from mininet.clean import cleanup

def bwtest( cpuLimits, period_us=100000, seconds=10 ):
    """Example/test of link and CPU bandwidth limits
    cpu: cpu limit as fraction of overall CPU time"""

    topo = TreeTopo( depth=1, fanout=2 )

    results = {}

    for sched in 'rt', 'cfs':
        info( '*** Testing with', sched, 'bandwidth limiting\n' )
```

^G Get Help
^X Exit

^O Write Out
^R Read File

^W Where Is
^_ Replace

^K Cut Text
^U Uncut Text

^J Justify
^T To Linter

^C Cur Pos
^_ Go To Line

M-U Undo
M-E Redo

```
from mininet.node import CPULimitedHost
from mininet.topolib import TreeTopo
from mininet.util import custom, waitListening, decode
from mininet.log import setLogLevel, info
from mininet.clean import cleanup

def bwtest( cpuLimits, period_us=100000, seconds=10 ):
    """Example/test of link and CPU bandwidth limits
    cpu: cpu limit as fraction of overall CPU time"""

    topo = TreeTopo( depth=1, fanout=2 )

    results = {}

    for sched in 'rt', 'cfs':
        info( '*** Testing with', sched, 'bandwidth limiting\n' )
        for cpu in cpuLimits:
            # cpu is the cpu fraction for all hosts, so we divide
            # it across two hosts
            host = custom( CPULimitedHost, sched=sched,
                           period_us=period_us,
                           cpu=.5*cpu )

            try:
                net = Mininet( topo=topo, host=host, waitConnected=True )
                # pylint: disable=bare-except
            except: # noqa
                info( '*** Skipping scheduler %s and cleaning up\n' % sched )
                cleanup()
                break

            net.start()
            net.pingAll()
            hosts = [ net.getNodeByName( h ) for h in topo.hosts() ]
```

^G Get Help
^X Exit

^O Write Out
^R Read File

^W Where Is
^_ Replace

^K Cut Text
^U Uncut Text

^J Justify
^T To Linter

^C Cur Pos
^_ Go To Line

M-U Undo
M-E Redo

```
for sched in 'rt', 'cfs':
    info( '*** Testing with', sched, 'bandwidth limiting\n' )
    for cpu in cpuLimits:
        # cpu is the cpu fraction for all hosts, so we divide
        # it across two hosts
        host = custom( CPULimitedHost, sched=sched,
                       period_us=period_us,
                       cpu=.5*cpu )

        try:
            net = Mininet( topo=topo, host=host, waitConnected=True )
            # pylint: disable=bare-except
        except: # noqa
            info( '*** Skipping scheduler %s and cleaning up\n' % sched )
            cleanup()
            break

        net.start()
        net.pingAll()
        hosts = [ net.getNodeByName( h ) for h in topo.hosts() ]
        client, server = hosts[ 0 ], hosts[ -1 ]
        info( '*** Starting iperf with %d%% of CPU allocated to hosts\n' %
              ( 100.0 * cpu ) )
        # We measure at the server because it doesn't include
        # the client's buffer fill rate
        popen = server.popen( 'iperf -yc -s -p 5001' )
        waitListening( client, server, 5001 )
        # ignore empty result from waitListening/telnet
        popen.stdout.readline()
        client.cmd( 'iperf -yc -t %s -c %s' % ( seconds, server.IP() ) )
        result = decode( popen.stdout.readline() ).split( ',' )
        bps = float( result[ -1 ] )
        popen.terminate()
        net.stop()
```

^G Get Help
^X Exit

^O Write Out
^R Read File

^W Where Is
^_ Replace

^K Cut Text
^U Uncut Text

^J Justify
^T To Linter

^C Cur Pos
^_ Go To Line

M-U Undo
M-E Redo

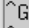

```
net.start()
net.pingAll()
hosts = [ net.getNodeByName( h ) for h in topo.hosts() ]
client, server = hosts[ 0 ], hosts[ -1 ]
info( '*** Starting iperf with %d%% of CPU allocated to hosts\n' %
      ( 100.0 * cpu ) )
# We measure at the server because it doesn't include
# the client's buffer fill rate
popen = server.popen( 'iperf -yc -s -p 5001' )
waitListening( client, server, 5001 )
# ignore empty result from waitListening/telnet
popen.stdout.readline()
client.cmd( 'iperf -yc -t %s -c %s' % ( seconds, server.IP() ) )
result = decode( popen.stdout.readline() ).split( ',' )
bps = float( result[ -1 ] )
popen.terminate()
net.stop()
updated = results.get( sched, [] )
updated += [ ( cpu, bps ) ]
results[ sched ] = updated



return results



def dump( results ):
    "Dump results"



    fmt = '%s\t%s\t%s\n'



    info( '\n' )
    info( fmt % ( 'sched', 'cpu', 'received bits/sec' ) )
```



 Get Help
 Exit

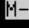
 Write Out
 Read File

 Where Is
 Replace

 Cut Text
 Uncut Text

 Justify
 To Linter

 Cur Pos
 Go To Line

 Undo
 Redo


```
GNU nano 2.9.3                                cpu.py

    net.stop()
    updated = results.get( sched, [] )
    updated += [ ( cpu, bps ) ]
    results[ sched ] = updated

    return results

def dump( results ):
    "Dump results"

    fmt = '%s\t%s\t%s\n'

    info( '\n' )
    info( fmt % ( 'sched', 'cpu', 'received bits/sec' ) )

    for sched in sorted( results.keys() ):
        entries = results[ sched ]
        for cpu, bps in entries:
            pct = '%d%%' % ( cpu * 100 )
            mbps = '%.2e' % bps
            info( fmt % ( sched, pct, mbps ) )

if __name__ == '__main__':
    setLogLevel( 'info' )
    # These are the limits for the hosts/iperfs - the
    # rest is for system processes
    limits = [ .5, .4, .3, .2, .1 ]
    out = bwtest( limits )
    dump( out )

^G Get Help      ^O Write Out    ^W Where Is     ^K Cut Text     ^J Justify      ^C Cur Pos      M-U Undo
^X Exit          ^R Read File    ^_ Replace      ^U Uncut Text   ^I To Linter    ^_ Go To Line    M-E Redo
```

CODE EXPLANATION:

- The first statement is a string literal, this is the module documentation string, explaining what the app does, with the help of docstring tools, user can interactively browse through the code.
- The six import statements are used to import **Mininet**, **CPULimitedHost**, **TreeTopo**, **custom**, **waitListening**, **decode**, **setLogLevel**, **info**, and **cleanup** modules from their parent modules respectively.
- **def** statement is used to define bwtest function that instantiates TreeTopo class with local variables depth = 1, and fanout = 2, assigned an empty dictionary to variable results
 - Used nested **for** loop to iterates two strings 'rt', and 'cfs'
 - Used **try** statement to handle exception within the inner for loop that iterates cpuLimits, called a custom functions with for arguments where cpu is divided by 2 (cpu=.5cpu)
 - The try clause is executed:
 - If no exception occurs, the except clause is skipped, and the code continues

- If exception occurs, the except clause is executed, and finished
 - **net.start()** start controller(s), switch(es)
 - **net.pingAll()** ping all hosts
 - A list of host names is assigned to hosts
 - **hosts[0]** is assigned to client, and **hosts[-1]** is assigned server
 - Used **iperf** to measure at the server and assigned to popen
 - Called client **cmd()** method with **iperf** command as argument
 - Formated the **popen** standard output
 - Returned **results**
- **def** statement is used to define a dump function that is used to dump the results return in the bwtest function.
 - The **if __name__ == "__main__":** statement makes the file usable as a script and importable module.

RESULT:

```
File Machine View Input Devices Help
ks-sdn@sdn_server:~/mininet/examples$ sudo python cpu.py
```

```
File Machine View Input Devices Help
ps ax | egrep -o 'dp[0-9]+' | sed 's/dp/nl:/'
*** Removing OVS datapaths
ovs-vsctl --timeout=1 list-br
ovs-vsctl --timeout=1 list-br
*** Removing all links of the pattern foo-ethX
ip link show | egrep -o '([-_.[:alnum:]]+-eth[[:digit:]]+)'
ip link show
*** Killing stale mininet node processes
pkill -9 -f mininet:
*** Shutting down stale tunnels
pkill -9 -f Tunnel=Ethernet
pkill -9 -f .ssh/mn
rm -f ~/.ssh/mn/*
*** Cleanup complete.
*** Testing with cfs bandwidth limiting
*** Creating network
*** Adding controller
*** Adding hosts:
h1 h2
*** Adding switches:
s1
*** Adding links:
(s1, h1) (s1, h2)
*** Configuring hosts
h1 (cfs 25000/100000us) h2 (cfs 25000/100000us)
*** Starting controller
c0
*** Starting 1 switches
s1 ...
*** Waiting for switches to connect
s1
*** Ping: testing ping reachability
h1 -> h2
h2 -> h1
*** Results: 0% dropped (2/2 received)
*** Starting iperf with 50% of CPU allocated to hosts
```

```
File Machine View Input Devices Help
h2 -> h1
*** Results: 0% dropped (2/2 received)
*** Starting iperf with 40% of CPU allocated to hosts
*** Stopping 1 controllers
c0
*** Stopping 2 links
...
*** Stopping 1 switches
s1
*** Stopping 2 hosts
h1 h2
*** Done
*** Creating network
*** Adding controller
*** Adding hosts:
h1 h2
*** Adding switches:
s1
*** Adding links:
(s1, h1) (s1, h2)
*** Configuring hosts
h1 (cfs 15000/1000000us) h2 (cfs 15000/1000000us)
*** Starting controller
c0
*** Starting 1 switches
s1 ...
*** Waiting for switches to connect
s1
*** Ping: testing ping reachability
h1 -> h2
h2 -> h1
*** Results: 0% dropped (2/2 received)
*** Starting iperf with 30% of CPU allocated to hosts
*** Stopping 1 controllers
c0
*** Stopping 2 links
..
```

```
File  Machine  View  Input  Devices  Help
s1
*** Ping: testing ping reachability
h1 -> h2
h2 -> h1
*** Results: 0% dropped (2/2 received)
*** Starting iperf with 20% of CPU allocated to hosts
*** Stopping 1 controllers
c0
*** Stopping 2 links
..
*** Stopping 1 switches
s1
*** Stopping 2 hosts
h1 h2
*** Done
*** Creating network
*** Adding controller
*** Adding hosts:
h1 h2
*** Adding switches:
s1
*** Adding links:
(s1, h1) (s1, h2)
*** Configuring hosts
h1 (cfs 5000/1000000us) h2 (cfs 5000/1000000us)
*** Starting controller
c0
*** Starting 1 switches
s1 ...
*** Waiting for switches to connect
s1
*** Ping: testing ping reachability
h1 -> h2
h2 -> h1
*** Results: 0% dropped (2/2 received)
*** Starting iperf with 10% of CPU allocated to hosts
```

```

File  Machine  View  Input  Devices  Help
*** Adding controller
*** Adding hosts:
h1 h2
*** Adding switches:
s1
*** Adding links:
(s1, h1) (s1, h2)
*** Configuring hosts
h1 (cfs 5000/1000000us) h2 (cfs 5000/1000000us)
*** Starting controller
c0
*** Starting 1 switches
s1 ...
*** Waiting for switches to connect
s1
*** Ping: testing ping reachability
h1 -> h2
h2 -> h1
*** Results: 0% dropped (2/2 received)
*** Starting iperf with 10% of CPU allocated to hosts
*** Stopping 1 controllers
c0
*** Stopping 2 links
..
*** Stopping 1 switches
s1
*** Stopping 2 hosts
h1 h2
*** Done

sched  cpu      received bits/sec
cfs    50%      8.94e+09
cfs    40%      7.00e+09
cfs    30%      4.22e+09
cfs    20%      3.15e+09
cfs    10%      1.50e+09
ks-sdn@sdn_server:~/mininet/examples$ _

```

Summary: The API display the cpu usage and the bandwidth.

API - 4

```
File Machine View Input Devices Help
GNU nano 2.9.3 emptynet.py

#!/usr/bin/env python
"""
This example shows how to create an empty Mininet object
(without a topology object) and add nodes to it manually.
"""

from mininet.net import Mininet
from mininet.node import Controller
from mininet.cli import CLI
from mininet.log import setLogLevel, info

def emptyNet():

    "Create an empty network and add nodes to it."

    net = Mininet( controller=Controller, waitConnected=True )

    info( '*** Adding controller\n' )
    net.addController( 'c0' )

    info( '*** Adding hosts\n' )
    h1 = net.addHost( 'h1', ip='10.0.0.1' )
    h2 = net.addHost( 'h2', ip='10.0.0.2' )

    info( '*** Adding switch\n' )
    s3 = net.addSwitch( 's3' )

    info( '*** Creating links\n' )
    net.addLink( h1, s3 )
    net.addLink( h2, s3 )

[ Read 45 lines ]
^G Get Help      ^O Write Out    ^W Where Is     ^K Cut Text     ^J Justify      ^C Cur Pos      M-U Undo
^X Exit          ^R Read File    ^_ Replace      ^U Uncut Text   ^T To Linter    ^_ Go To Line    M-E Redo
```

```
File Machine View Input Devices Help
GNU nano 2.9.3 emptynet.py

net = Mininet( controller=Controller, waitConnected=True )

info( '*** Adding controller\n' )
net.addController( 'c0' )

info( '*** Adding hosts\n' )
h1 = net.addHost( 'h1', ip='10.0.0.1' )
h2 = net.addHost( 'h2', ip='10.0.0.2' )

info( '*** Adding switch\n' )
s3 = net.addSwitch( 's3' )

info( '*** Creating links\n' )
net.addLink( h1, s3 )
net.addLink( h2, s3 )

info( '*** Starting network\n' )
net.start()

info( '*** Running CLI\n' )
CLI( net )

info( '*** Stopping network\n' )
net.stop()

if __name__ == '__main__':
    setLogLevel( 'info' )
    emptyNet()
```

Get Help Write Out Where Is Cut Text Justify Cur Pos M-U Undo
Exit Read File Replace Uncut Text To Linter Go To Line M-E Redo

CODE EXPLANATION

- The first statement is a string literal, this is the module documentation string, explaining what the app does, with the help of docstring tools, user can interactively browse through the code.
- The four import statements are used to **import Mininet, Controller, setLogLevel, Info, and CLI modules** from their parent modules respectively.
- The def statement defined a function named emptyNet which was used to:
 - instance of **Mininet** class was instantiated and assigned to a variable **net**
 - Reference **addController** method with the instance object **net**,
 - Reference **addHost** method with the instance object **net** to add hosts **h1, h2**
 - Reference **addSwitch** method with the instance object **net** to add switch **s3**
 - Reference **addLink** method with the instance object **net** to add links (**h1 - s3, h2 - s3**)
 - Start the controller, switch, and hosts with **net.start()**
 - Stop controller, switch, and hosts with **net.stop()**
 -

- The `if __name__ == "__main__":` statement makes the file usable as a script and importable module.

RESULT:

```

File  Machine  View  Input  Devices  Help
ks-sdn@sdn_server:~$ cd mininet/examples/
ks-sdn@sdn_server:~/mininet/examples$ sudo python emptynet.py
[sudo] password for ks-sdn:
*** Adding controller
*** Adding hosts
*** Adding switch
*** Creating links
*** Starting network
*** Configuring hosts
h1 h2
*** Starting controller
c0
*** Starting 1 switches
s3 ...
*** Waiting for switches to connect
s3
*** Running CLI
*** Starting CLI:
mininet> pingall
*** Ping: testing ping reachability
h1 -> h2
h2 -> h1
*** Results: 0% dropped (2/2 received)
mininet> net
h1 h1-eth0:s3-eth1
h2 h2-eth0:s3-eth2
s3 lo: s3-eth1:h1-eth0 s3-eth2:h2-eth0
c0
mininet> dump
<Host h1: h1-eth0:10.0.0.1 pid=2294>
<Host h2: h2-eth0:10.0.0.2 pid=2296>
<OVSSwitch s3: lo:127.0.0.1,s3-eth1:None,s3-eth2:None pid=2301>
<Controller c0: 127.0.0.1:6653 pid=2287>
mininet>

```

APP - 5

```
#!/usr/bin/env python
```

```
"""
```

```
This example shows how to add an interface (for example a real  
hardware interface) to a network after the network is created.
```

```
"""
```

```
import re  
import sys
```

```
from sys import exit # pylint: disable=redefined-builtin
```

```
from mininet.cli import CLI  
from mininet.log import setLogLevel, info, error  
from mininet.net import Mininet  
from mininet.link import Intf  
from mininet.topolib import TreeTopo  
from mininet.util import quietRun
```

```
def checkIntf( intf ):  
    "Make sure intf exists and is not configured."  
    config = quietRun( 'ifconfig %s 2>/dev/null' % intf, shell=True )  
    if not config:  
        error( 'Error:', intf, 'does not exist!\n' )  
        exit( 1 )  
    ips = re.findall( r'\d+\.\d+\.\d+\.\d+', config )  
    if ips:  
        error( 'Error:', intf, 'has an IP address,'  
              'and is probably in use!\n' )  
        exit( 1 )
```

```
^G Get Help  
^X Exit
```

```
^O Write Out  
^R Read File
```

```
^W Where Is  
^_ Replace
```

```
^K Cut Text  
^U Uncut Text
```

```
^J Justify  
^T To Linter
```

```
^C Cur Pos  
^_ Go To Line
```

```
M-U Undo  
M-E Redo
```

```
from mininet.link import Intf
from mininet.topolib import TreeTopo
from mininet.util import quietRun

def checkIntf( intf ):
    "Make sure intf exists and is not configured."
    config = quietRun( 'ifconfig %s 2>/dev/null' % intf, shell=True )
    if not config:
        error( 'Error:', intf, 'does not exist!\n' )
        exit( 1 )
    ips = re.findall( r'\d+\.\d+\.\d+\.\d+', config )
    if ips:
        error( 'Error:', intf, 'has an IP address,'
              'and is probably in use!\n' )
        exit( 1 )

if __name__ == '__main__':
    setLogLevel( 'info' )

    # try to get hw intf from the command line; by default, use eth1
    intfName = sys.argv[ 1 ] if len( sys.argv ) > 1 else 'eth1'
    info( '*** Connecting to hw intf: %s' % intfName )

    info( '*** Checking', intfName, '\n' )
    checkIntf( intfName )

    info( '*** Creating network\n' )
    net = Mininet( topo=TreeTopo( depth=1, fanout=2 ), waitConnected=True )

    switch = net.switches[ 0 ]
```

^G Get Help
^X Exit

^O Write Out
^R Read File

^W Where Is
^_ Replace

^K Cut Text
^U Uncut Text

^J Justify
^T To Linter

^C Cur Pos
^_ Go To Line

M-U Undo
M-E Redo

```
GNU nano 2.9.3                               hwintf.py

if __name__ == '__main__':
    setLogLevel( 'info' )

    # try to get hw intf from the command line; by default, use eth1
    intfName = sys.argv[ 1 ] if len( sys.argv ) > 1 else 'eth1'
    info( '*** Connecting to hw intf: %s' % intfName )

    info( '*** Checking', intfName, '\n' )
    checkIntf( intfName )

    info( '*** Creating network\n' )
    net = Mininet( topo=TreeTopo( depth=1, fanout=2 ), waitConnected=True )

    switch = net.switches[ 0 ]
    info( '*** Adding hardware interface', intfName, 'to switch',
          switch.name, '\n' )
    _intf = Intf( intfName, node=switch )

    info( '*** Note: you may need to reconfigure the interfaces for '
          'the Mininet hosts:\n', net.hosts, '\n' )

    net.start()
    CLI( net )
    net.stop()
```

^G Get Help ^O Write Out ^W Where Is ^K Cut Text ^J Justify ^C Cur Pos M-U Undo
^X Exit ^R Read File ^\ Replace ^U Uncut Text ^T To Linter ^_ Go To Line M-E Redo

CODE EXPLANATION

- The first statement is a string literal, this is the module documentation string, explaining what the app does, with the help of docstring tools, user can interactively browse through the code.
- The nine **import** statements are used to import **re**, **sys**, **CLI**, **setLogLevel**, **Info**, **error**, **Mininet**, **Intf**, **TreeTopo**, **quietRun** from their parent modules respectively.
- The **def** statement defined **checkIntf** method, that check if an interface exist and configured by:
 - Calling a **quietRun** method that performed an **ifconfig** command, redirect the **stderr** to a null device, and assigned to variable **config**
 - Used **if not** statement to verify **config** which raised an **exception** and exit with **code 1**: operation not permitted.
 - Called the **findall** method and assigned return value to **ips**
 - Used **if** statement to check **ips**, raised **exception** and exit with **code 1**, operation not permitted.

- The `if __name__ == "__main__":` statement makes the file usable as a script and importable module.
 - Used `sys.argv` to accept input from the command line, set the default to `eth1`, assigned to variable `intfName`
 - Called `checkIntf` method with `intfName` as argument
 - Instantiates `Mininet` class and assigned it to variable `net`
 - Reference the `Mininet` class method with `net.switches[0]`
 - Start mininet network
 - **CLI(net)** runs a batch or interactive mode for the network.
 - The `net.stop()` stop controllers, switches, and hosts.

```
File Machine View Input Devices Help
ks-sdn@sdn_server:~/mininet/examples$ sudo python hwintf.py
*** Connecting to hw intf: eth1*** Checking eth1
Error: eth1 does not exist!
ks-sdn@sdn_server:~/mininet/examples$ sudo python hwintf.py 192.168.1.1
*** Connecting to hw intf: 192.168.1.1*** Checking 192.168.1.1
Error: 192.168.1.1 does not exist!
ks-sdn@sdn_server:~/mininet/examples$ sudo python hwintf.py eth2
*** Connecting to hw intf: eth2*** Checking eth2
Error: eth2 does not exist!
ks-sdn@sdn_server:~/mininet/examples$
```

Overlays: OpenVswitch + VXLAN

- All required software already install in iteration 1
- Create a directory named **overlay-lab** where all files for the project will be kept
- Create a **.yml file**, name it **docker-compose.yml** with the vi editor as shown below:

```
File Machine View Input Devices Help
ks-sdn@sdn_server:~$ cd overlay-lab/
ks-sdn@sdn_server:~/overlay-lab$ vim docker-compose.yml _
```



```

File Machine View Input Devices Help
ks-sdn@sdn_server:~$ cd overlay-lab/
ks-sdn@sdn_server:~/overlay-lab$ ls
docker-compose.yml  vm1_overlay.pcap  vm2_overlay.pcap
start_ovs.sh        vm1_underlay.pcap vm2_underlay.pcap
ks-sdn@sdn_server:~/overlay-lab$ sudo docker-compose up -d
/usr/lib/python2.7/dist-packages/OpenSSL/crypto.py:12: CryptographyDeprecationWarning: Python 2 is no longer supported by the Python core team. Support for it is now deprecated in cryptography, and will be removed in the next release.
  from cryptography import x509
overlaylab_hv1_1 is up-to-date
overlaylab_hv2_1 is up-to-date
ks-sdn@sdn_server:~/overlay-lab$ sudo docker ps

```

| CONTAINER ID | IMAGE | COMMAND | CREATED | STATUS |
|--------------|---------------------------------|--------------------------|------------|------------|
| 7c55d5215f37 | hechaol/ovs-vtep-emulator:2.9.0 | "/bin/bash -c '/star..." | 3 days ago | Up 2 hours |
| 1f86c3a5a519 | hechaol/ovs-vtep-emulator:2.9.0 | "/bin/bash -c '/star..." | 3 days ago | Up 2 hours |

```

ks-sdn@sdn_server:~/overlay-lab$ _

```

Configure Container 1, bridge and interfaces, local interface, and VTEP emulator:


```

File  Machine  View  Input  Devices  Help
onds          overlaylab_hv2_1
ks-sdn@sdn_server:~/overlay-lab$ sudo docker exec -it overlaylab_hv1_1 /bin/bash
root@7c55d5215f37:/# ovs-vsctl add-br br0
root@7c55d5215f37:/# ovs-vsctl add-port br0 vm1 --set interface vm1 type=internal
ovs-vsctl: Port does not contain a column whose name matches "--set"
root@7c55d5215f37:/# ovs-vsctl add-port br0 vm1 -- set interface vm1 type=internal
root@7c55d5215f37:/# ifconfig vm1 hw ether 00:00:00:11:11:11
root@7c55d5215f37:/# ifconfig vm1 192.168.1.1 netmask 255.255.255.240 up
root@7c55d5215f37:/#
root@7c55d5215f37:/# vtep-ctl add-ps br0
root@7c55d5215f37:/# vtep-ctl set Physical_SW br0 tunnel_ips=10.0.0.1
root@7c55d5215f37:/# /usr/local/share/openvswitch/scripts/ovs-vtep --log-file --pidfile --detach br0
root@7c55d5215f37:/# vtep-ctl add-ls vswitch0
root@7c55d5215f37:/# vtep-ctl set Logical_SW vswitch0 tunnel_key=5000
root@7c55d5215f37:/# vtep-ctl bind-ls br0 vm1 0 vswitch0
root@7c55d5215f37:/# vtep-ctl add-ucast-remote vswitch0
vtep-ctl: 'add-ucast-remote' command requires at least 3 arguments
root@7c55d5215f37:/# vtep-ctl add-ucast-remote vswitch0 00:00:00:22:22:22 10.1.0.1
root@7c55d5215f37:/# vtep-ctl add-mcast-remote vswitch0 unknown-dst 10.1.0.1
root@7c55d5215f37:/# vtep-ctl list Physical_SW
 _uuid      : 1ccf32e0-f46e-430c-b90b-b470d874c74f
description : "OVS VTEP Emulator"
management_ips : []
name        : "br0"
other_config : {}
ports       : [664523ee-b3b3-46af-9499-cdae46bcb092]
switch_fault_status : []
tunnel_ips  : ["10.0.0.1"]
tunnels     : [7bb4fc9c-dde6-42d5-b6fd-45c69b2c30ae]
root@7c55d5215f37:/# vtep-ctl list Logical_SW
 _uuid      : 9d5789f0-6574-400b-bd13-c49c661b995c
description : ""
name        : "vswitch0"
other_config : {}
replication_mode : []
tunnel_key  : 5000
root@7c55d5215f37:/#

```

Configure Container 2, bridge and interfaces, local interface, and VTEP emulator.

| File | Machine | View | Input | Devices | Help |
|-----------------------------------|---------------------------------------------------------------------------------|--------------------------|----------------|-----------|------|
| utes | overlaylab_hv1_1 | | | | |
| 1f86c3a5a519 | hechaol/ovs-vtep-emulator:2.9.0 | "/bin/bash -c '/star..." | 33 minutes ago | Up 33 min | |
| utes | overlaylab_hv2_1 | | | | |
| ks-sdn@sdn_server:~/overlay-lab\$ | sudo docker exec -it overlay_hv2_1 /bin/bash | | | | |
| Error: | No such container: overlay_hv2_1 | | | | |
| ks-sdn@sdn_server:~/overlay-lab\$ | sudo docker exec -it overlaylab_hv2_1 /bin/bash | | | | |
| root@1f86c3a5a519:/# | ovs-vsctl add-br br0 | | | | |
| root@1f86c3a5a519:/# | ovs-vsctl add-port br0 vm2 -- set interface vm2 type=internal | | | | |
| root@1f86c3a5a519:/# | ifconfig vm2 hw ether 00:00:00:22:22:22 | | | | |
| root@1f86c3a5a519:/# | ifconfig vm2 192.168.1.2 netmask 255.255.255.240 up | | | | |
| root@1f86c3a5a519:/# | | | | | |
| root@1f86c3a5a519:/# | vtep-ctl add-ps br0 | | | | |
| root@1f86c3a5a519:/# | vtep-ctl set Physical_SW br0 tunnel_ips=10.1.0.1 | | | | |
| root@1f86c3a5a519:/# | /usr/local/share/openvswitch/scripts/ovs-vtep --log-file --pidfile --detach br0 | | | | |
| root@1f86c3a5a519:/# | vtep-ctl add-ls vswitch0 | | | | |
| root@1f86c3a5a519:/# | vtep-ctl set Logical_SW vswitch0 tunnel_key=5000 | | | | |
| root@1f86c3a5a519:/# | vtep-ctl bind-ls br0 vm2 0 vswitch0 | | | | |
| root@1f86c3a5a519:/# | vtep-ctl add-ucast-remote vswitch0 00:00:00:11:11:11 10.0.0.1 | | | | |
| root@1f86c3a5a519:/# | | | | | |
| root@1f86c3a5a519:/# | vtep-ctl list Physical_SW | | | | |
| _uuid | : 42ee6741-6b5a-46b4-bc93-c886b8038a8a | | | | |
| description | : "OVS VTEP Emulator" | | | | |
| management_ips | : [] | | | | |
| name | : "br0" | | | | |
| other_config | : {} | | | | |
| ports | : [85cb257d-a5c8-4484-a51c-c81eafe054f3] | | | | |
| switch_fault_status | : [] | | | | |
| tunnel_ips | : ["10.1.0.1"] | | | | |
| tunnels | : [6b9db6f9-15cb-4af9-9f5b-7eaf3e041bb9] | | | | |
| root@1f86c3a5a519:/# | vtep-ctl list Logical_SW | | | | |
| _uuid | : 7c7cbf97-f84d-4db8-ad25-9fb061679f4d | | | | |
| description | : "" | | | | |
| name | : "vswitch0" | | | | |
| other_config | : {} | | | | |
| replication_mode | : [] | | | | |
| tunnel_key | : 5000 | | | | |
| root@1f86c3a5a519:/# | _ | | | | |

Testing Configurations:

```

File Machine View Input Devices Help
ks-sdn@sdn_server:~/overlay-lab$ sudo docker ps
CONTAINER ID        IMAGE                                     COMMAND                  CREATED            STATUS
7c55d5215f37       hechaol/ovs-vtep-emulator:2.9.0        "/bin/bash -c '/star..." About an hour ago   Up Abo
ut an hour         overlaylab_hv1_1
1f86c3a5a519       hechaol/ovs-vtep-emulator:2.9.0        "/bin/bash -c '/star..." About an hour ago   Up Abo
ut an hour         overlaylab_hv2_1
ks-sdn@sdn_server:~/overlay-lab$
ks-sdn@sdn_server:~/overlay-lab$ sudo docker exec overlaylab_hv1_1 ping -I vm1 192.168.1.2 -c1
PING 192.168.1.2 (192.168.1.2) from 192.168.1.1 vm1: 56(84) bytes of data.
64 bytes from 192.168.1.2: icmp_seq=1 ttl=64 time=1.54 ms

--- 192.168.1.2 ping statistics ---
1 packets transmitted, 1 received, 0% packet loss, time 0ms
rtt min/avg/max/mdev = 1.541/1.541/1.541/0.000 ms
ks-sdn@sdn_server:~/overlay-lab$
ks-sdn@sdn_server:~/overlay-lab$ sudo docker exec overlaylab_hv2_1 ping -I vm2 192.168.1.1 -c1
PING 192.168.1.1 (192.168.1.1) from 192.168.1.2 vm2: 56(84) bytes of data.
64 bytes from 192.168.1.1: icmp_seq=1 ttl=64 time=0.882 ms

--- 192.168.1.1 ping statistics ---
1 packets transmitted, 1 received, 0% packet loss, time 0ms
rtt min/avg/max/mdev = 0.882/0.882/0.882/0.000 ms
ks-sdn@sdn_server:~/overlay-lab$
ks-sdn@sdn_server:~/overlay-lab$

```

Sniff traffics using tcpdump command:

```

File Machine View Input Devices Help
ks-sdn@sdn_server:~/overlay-lab$
ks-sdn@sdn_server:~/overlay-lab$ sudo docker ps
CONTAINER ID        IMAGE                                     COMMAND                  CREATED            STATUS
PORTS              NAMES
7c55d5215f37       hechaol/ovs-vtep-emulator:2.9.0       "/bin/bash -c '/star..." 3 days ago        Up 6 minutes
overlaylab_hv1_1
1f86c3a5a519       hechaol/ovs-vtep-emulator:2.9.0       "/bin/bash -c '/star..." 3 days ago        Up 6 minutes
overlaylab_hv2_1
ks-sdn@sdn_server:~/overlay-lab$ sudo docker exec -it overlaylab_hv1_1 /bin/bash
root@7c55d5215f37:/# ifconfig vm1 hw ether 00:00:00:11:11:11
root@7c55d5215f37:/# ifconfig vm1 192.168.1.1 netmask 255.255.255.240 up
root@7c55d5215f37:/# vtep-ctl list Physical_SW
_uuid                : 1ccf32e0-f46e-430c-b90b-b470d874c74f
description          : "OVS VTEP Emulator"
management_ips       : []
name                 : "br0"
other_config          : {}
ports                : [664523ee-b3b3-46af-9499-cdae46bcb092]
switch_fault_status  : []
tunnel_ips            : ["10.0.0.1"]
tunnels               : [7bb4fc9c-dde6-42d5-b6fd-45c69b2c30ae]
root@7c55d5215f37:/# vtep-ctl list Logical_SW
_uuid                : 9d5789f0-6574-400b-bd13-c49c661b995c
description          : ""
name                 : "vswitch0"
other_config          : {}
replication_mode      : []
tunnel_key            : 5000
root@7c55d5215f37:/# tcpdump -i eth0 -w vm1_underlay.pcap
tcpdump: listening on eth0, link-type EN10MB (Ethernet), capture size 262144 bytes
^C33 packets captured
33 packets received by filter
0 packets dropped by kernel
root@7c55d5215f37:/# _

```

```
File Machine View Input Devices Help
ks-sdn@sdn_server:~$ sudo docker exec -it overlaylab_hv2_1 /bin/bash
[sudo] password for ks-sdn:
root@1f86c3a5a519:/# ifconfig vm2 hw ether 00:00:00:22:22:22
root@1f86c3a5a519:/# ifconfig vm2 192.168.1.2 netmask 255.255.255.240 up
root@1f86c3a5a519:/# vtep-ctl list Physical_SW
 _uuid      : 42ee6741-6b5a-46b4-bc93-c886b8038a8a
description : "OVS VTEP Emulator"
management_ips : []
name        : "br0"
other_config : {}
ports       : [85cb257d-a5c8-4484-a51c-c81eafe054f3]
switch_fault_status : []
tunnel_ips  : ["10.1.0.1"]
tunnels     : [6b9db6f9-15cb-4af9-9f5b-7eaf3e041bb9]
root@1f86c3a5a519:/# vtep-ctl list Logical_SW
 _uuid      : 7c7cbf97-f84d-4db8-ad25-9fb061679f4d
description : ""
name        : "vswitch0"
other_config : {}
replication_mode : []
tunnel_key   : 5000
root@1f86c3a5a519:/# tcpdump -i vm2 -w vm2_overlay.pcap
tcpdump: listening on vm2, link-type EN10MB (Ethernet), capture size 262144 bytes
^C26 packets captured
26 packets received by filter
0 packets dropped by kernel
root@1f86c3a5a519:/# tcpdump -i eth0 -w vm2_underlay.pcap
tcpdump: listening on eth0, link-type EN10MB (Ethernet), capture size 262144 bytes
^C14 packets captured
14 packets received by filter
0 packets dropped by kernel
root@1f86c3a5a519:/# _
```

```

File Machine View Input Devices Help
ks-sdn@sdn_server:~/overlay-lab$
ks-sdn@sdn_server:~/overlay-lab$ sudo docker exec overlaylab_hv1_1 ping -I vm1 192.168.1.2 -c 5
PING 192.168.1.2 (192.168.1.2) from 192.168.1.1 vm1: 56(84) bytes of data.
64 bytes from 192.168.1.2: icmp_seq=1 ttl=64 time=2.61 ms
64 bytes from 192.168.1.2: icmp_seq=2 ttl=64 time=0.195 ms
64 bytes from 192.168.1.2: icmp_seq=3 ttl=64 time=0.199 ms
64 bytes from 192.168.1.2: icmp_seq=4 ttl=64 time=0.169 ms
64 bytes from 192.168.1.2: icmp_seq=5 ttl=64 time=0.230 ms

--- 192.168.1.2 ping statistics ---
5 packets transmitted, 5 received, 0% packet loss, time 4056ms
rtt min/avg/max/mdev = 0.169/0.682/2.618/0.968 ms
ks-sdn@sdn_server:~/overlay-lab$ sudo docker exec overlaylab_hv2_1 ping -I vm2 192.168.1.1 -c 5
PING 192.168.1.1 (192.168.1.1) from 192.168.1.2 vm2: 56(84) bytes of data.
64 bytes from 192.168.1.1: icmp_seq=1 ttl=64 time=1.12 ms
64 bytes from 192.168.1.1: icmp_seq=2 ttl=64 time=0.190 ms
64 bytes from 192.168.1.1: icmp_seq=3 ttl=64 time=0.187 ms
64 bytes from 192.168.1.1: icmp_seq=4 ttl=64 time=0.190 ms
64 bytes from 192.168.1.1: icmp_seq=5 ttl=64 time=0.222 ms

--- 192.168.1.1 ping statistics ---
5 packets transmitted, 5 received, 0% packet loss, time 4044ms
rtt min/avg/max/mdev = 0.187/0.383/1.129/0.373 ms
ks-sdn@sdn_server:~/overlay-lab$ sudo docker exec overlaylab_hv2_1 ping -I vm2 192.168.1.1 -c 5
PING 192.168.1.1 (192.168.1.1) from 192.168.1.2 vm2: 56(84) bytes of data.
64 bytes from 192.168.1.1: icmp_seq=1 ttl=64 time=0.907 ms
64 bytes from 192.168.1.1: icmp_seq=2 ttl=64 time=0.186 ms
64 bytes from 192.168.1.1: icmp_seq=3 ttl=64 time=0.186 ms
64 bytes from 192.168.1.1: icmp_seq=4 ttl=64 time=0.186 ms
64 bytes from 192.168.1.1: icmp_seq=5 ttl=64 time=0.247 ms

--- 192.168.1.1 ping statistics ---
5 packets transmitted, 5 received, 0% packet loss, time 4041ms
rtt min/avg/max/mdev = 0.186/0.342/0.907/0.283 ms
ks-sdn@sdn_server:~/overlay-lab$ _

```

Transfer Container Captured files to current directory:

```

File Machine View Input Devices Help
ks-sdn@sdn_server:~/overlay-lab$ sudo docker cp overlaylab_hv1_1:/vm1_overlay.pcap .
[sudo] password for ks-sdn:
ks-sdn@sdn_server:~/overlay-lab$ sudo docker cp overlaylab_hv2_1:/vm2_overlay.pcap .
ks-sdn@sdn_server:~/overlay-lab$ sudo docker cp overlaylab_hv1_1:/vm1_underlay.pcap .
ks-sdn@sdn_server:~/overlay-lab$ sudo docker cp overlaylab_hv2_1:/vm2_underlay.pcap .
ks-sdn@sdn_server:~/overlay-lab$ ls -l *.pcap
-rw-r--r-- 1 root root 2536 Jun 30 22:16 vm1_overlay.pcap
-rw-r--r-- 1 root root 4386 Jul  3 14:21 vm1_underlay.pcap
-rw-r--r-- 1 root root 2652 Jul  3 14:23 vm2_overlay.pcap
-rw-r--r-- 1 root root 1996 Jul  3 14:25 vm2_underlay.pcap
ks-sdn@sdn_server:~/overlay-lab$
ks-sdn@sdn_server:~/overlay-lab$ _

```

View files using tshark:

```
File Machine View Input Devices Help
ks-sdn@sdn_server:~/overlay-lab$ tshark -V -r vm1_overlay.pcap -c 1 --color_

File Machine View Input Devices Help
.... 0101 = Header Length: 20 bytes (5)
Differentiated Services Field: 0x00 (DSCP: CS0, ECN: Not-ECT)
  0000 00.. = Differentiated Services Codepoint: Default (0)
  .... ..00 = Explicit Congestion Notification: Not ECN-Capable Transport (0)
Total Length: 84
Identification: 0x1c71 (7281)
Flags: 0x4000, Don't fragment
  0... .... .... .... = Reserved bit: Not set
  .1.. .... .... .... = Don't fragment: Set
  ..0. .... .... .... = More fragments: Not set
  ...0 0000 0000 0000 = Fragment offset: 0
Time to live: 64
Protocol: ICMP (1)
Header checksum: 0x9ae4 [validation disabled]
[Header checksum status: Unverified]
Source: 192.168.1.1
Destination: 192.168.1.2
Internet Control Message Protocol
Type: 8 (Echo (ping) request)
Code: 0
Checksum: 0x1a2f [correct]
[Checksum Status: Good]
Identifier (BE): 82 (0x0052)
Identifier (LE): 20992 (0x5200)
Sequence number (BE): 1 (0x0001)
Sequence number (LE): 256 (0x0100)
Timestamp from icmp data: Jun 30, 2022 21:53:16.000000000 UTC
[Timestamp from icmp data (relative): 0.339323000 seconds]
Data (48 bytes)

0000  0f 2d 05 00 00 00 00 00 10 11 12 13 14 15 16 17  .-.....
0010  18 19 1a 1b 1c 1d 1e 1f 20 21 22 23 24 25 26 27  .....!#$%&'
0020  28 29 2a 2b 2c 2d 2e 2f 30 31 32 33 34 35 36 37  ()*+,-./01234567
      Data: 0f2d050000000000101112131415161718191a1b1c1d1e1f...
      [Length: 48]

ks-sdn@sdn_server:~/overlay-lab$
```

```
File Machine View Input Devices Help
ks-sdn@sdn_server:~/overlay-lab$ ls -l *.pcap
-rw-r--r-- 1 root root 2536 Jun 30 22:16 vm1_overlay.pcap
-rw-r--r-- 1 root root 4386 Jul  3 14:21 vm1_underlay.pcap
-rw-r--r-- 1 root root 2652 Jul  3 14:23 vm2_overlay.pcap
-rw-r--r-- 1 root root 1996 Jul  3 14:25 vm2_underlay.pcap
ks-sdn@sdn_server:~/overlay-lab$ tshark -V -r vm1_underlay.pcap -c 1 --color_

File Machine View Input Devices Help
ks-sdn@sdn_server:~/overlay-lab$
```

```

File Machine View Input Devices Help
[Coloring Rule String: eth[0] & 1]
Ethernet II, Src: 42:91:e9:ca:5d:f6 (42:91:e9:ca:5d:f6), Dst: IPv6mcast_02 (33:33:00:00:00:02)
  Destination: IPv6mcast_02 (33:33:00:00:00:02)
  Address: IPv6mcast_02 (33:33:00:00:00:02)
    .... ..1. .... = LG bit: Locally administered address (this is NOT the factory default)
    .... ..1. .... = IG bit: Group address (multicast/broadcast)
  Source: 42:91:e9:ca:5d:f6 (42:91:e9:ca:5d:f6)
  Address: 42:91:e9:ca:5d:f6 (42:91:e9:ca:5d:f6)
    .... ..1. .... = LG bit: Locally administered address (this is NOT the factory default)
    .... ..0. .... = IG bit: Individual address (unicast)
Type: IPv6 (0x86dd)
Internet Protocol Version 6, Src: fe80::4091:e9ff:feca:5df6, Dst: ff02::2
  0110 .... = Version: 6
  .... 0000 0000 .... = Traffic Class: 0x00 (DSCP: CS0, ECN: Not-ECT)
  .... 0000 00.. .... = Differentiated Services Codepoint: Default (0)
  .... .... ..00 .... = Explicit Congestion Notification: Not ECN-Capable Transport (0)
  .... .... 0000 0000 0000 0000 0000 = Flow Label: 0x000000
Payload Length: 16
Next Header: ICMPv6 (58)
Hop Limit: 255
Source: fe80::4091:e9ff:feca:5df6
Destination: ff02::2
Internet Control Message Protocol v6
Type: Router Solicitation (133)
Code: 0
Checksum: 0x6a89 [correct]
[Checksum Status: Good]
Reserved: 00000000
ICMPv6 Option (Source link-layer address : 42:91:e9:ca:5d:f6)
  Type: Source link-layer address (1)
  Length: 1 (8 bytes)
  Link-layer address: 42:91:e9:ca:5d:f6 (42:91:e9:ca:5d:f6)

ks-sdn@sdn_server:~/overlay-lab$ _

```

```

File Machine View Input Devices Help
ks-sdn@sdn_server:~/overlay-lab$ ls -l *.pcap
-rw-r--r-- 1 root root 2536 Jun 30 22:16 vm1_overlay.pcap
-rw-r--r-- 1 root root 4386 Jul  3 14:21 vm1_underlay.pcap
-rw-r--r-- 1 root root 2652 Jul  3 14:23 vm2_overlay.pcap
-rw-r--r-- 1 root root 1996 Jul  3 14:25 vm2_underlay.pcap
ks-sdn@sdn_server:~/overlay-lab$ tshark -V -r vm2_overlay.pcap -c 1 --color _

```



```

File  Machine  View  Input  Devices  Help
Arrival Time: Jul  3, 2022 14:17:02.499671000 UTC
[Time shift for this packet: 0.000000000 seconds]
Epoch Time: 1656857822.499671000 seconds
[Time delta from previous captured frame: 0.000000000 seconds]
[Time delta from previous displayed frame: 0.000000000 seconds]
[Time since reference or first frame: 0.000000000 seconds]
Frame Number: 1
Frame Length: 42 bytes (336 bits)
Capture Length: 42 bytes (336 bits)
[Frame is marked: False]
[Frame is ignored: False]
[Protocols in frame: eth:ethertype:arp]
[Coloring Rule Name: ARP]
[Coloring Rule String: arp]
Ethernet II, Src: 00:00:00_11:11:11 (00:00:00:11:11:11), Dst: Broadcast (ff:ff:ff:ff:ff:ff)
  Destination: Broadcast (ff:ff:ff:ff:ff:ff)
    Address: Broadcast (ff:ff:ff:ff:ff:ff)
      .... ..1. .... .... .... = LG bit: Locally administered address (this is NOT the factory default)
        .... ..1 .... .... .... = IG bit: Group address (multicast/broadcast)
    Source: 00:00:00_11:11:11 (00:00:00:11:11:11)
      Address: 00:00:00_11:11:11 (00:00:00:11:11:11)
        .... ..0. .... .... .... = LG bit: Globally unique address (factory default)
          .... ..0 .... .... .... = IG bit: Individual address (unicast)
    Type: ARP (0x0806)
Address Resolution Protocol (request)
  Hardware type: Ethernet (1)
  Protocol type: IPv4 (0x0800)
  Hardware size: 6
  Protocol size: 4
  Opcode: request (1)
  Sender MAC address: 00:00:00_11:11:11 (00:00:00:11:11:11)
  Sender IP address: 192.168.1.1
  Target MAC address: 00:00:00_00:00:00 (00:00:00:00:00:00)
  Target IP address: 192.168.1.2

ks-sdn@sdn_server:~/overlay-lab$

```

```

File  Machine  View  Input  Devices  Help
ks-sdn@sdn_server:~/overlay-lab$ ls -l *.pcap
-rw-r--r-- 1 root root 2536 Jun 30 22:16 vm1_overlay.pcap
-rw-r--r-- 1 root root 4386 Jul  3 14:21 vm1_underlay.pcap
-rw-r--r-- 1 root root 2652 Jul  3 14:23 vm2_overlay.pcap
-rw-r--r-- 1 root root 1996 Jul  3 14:25 vm2_underlay.pcap
ks-sdn@sdn_server:~/overlay-lab$ tshark -V -r vm2_underlay.pcap -c 1 --color

```

```

File  Machine  View  Input  Devices  Help
.... 0101 = Header Length: 20 bytes (5)
Differentiated Services Field: 0x00 (DSCP: CS0, ECN: Not-ECT)
    0000 00.. = Differentiated Services Codepoint: Default (0)
    .... ..00 = Explicit Congestion Notification: Not ECN-Capable Transport (0)
Total Length: 84
Identification: 0x18fe (6398)
Flags: 0x4000, Don't fragment
    0... .... = Reserved bit: Not set
    .1.. .... = Don't fragment: Set
    ..0. .... = More fragments: Not set
    ...0 0000 0000 0000 = Fragment offset: 0
Time to live: 64
Protocol: ICMP (1)
Header checksum: 0x9e57 [validation disabled]
[Header checksum status: Unverified]
Source: 192.168.1.2
Destination: 192.168.1.1
Internet Control Message Protocol
Type: 8 (Echo (ping) request)
Code: 0
Checksum: 0x037f [correct]
[Checksum Status: Good]
Identifier (BE): 46 (0x002e)
Identifier (LE): 11776 (0x2e00)
Sequence number (BE): 1 (0x0001)
Sequence number (LE): 256 (0x0100)
Timestamp from icmp data: Jul  3, 2022 14:25:21.000000000 UTC
[Timestamp from icmp data (relative): 0.554732000 seconds]
Data (48 bytes)

0000  9a 75 08 00 00 00 00 10 11 12 13 14 15 16 17  .U.....
0010  18 19 1a 1b 1c 1d 1e 1f 20 21 22 23 24 25 26 27  .....!'"$%&'
0020  28 29 2a 2b 2c 2d 2e 2f 30 31 32 33 34 35 36 37  ()*+,-./01234567
      Data: 9a7508000000000101112131415161718191a1b1c1d1e1f...
      [Length: 48]

ks-sdn@sdn_server:~/overlay-lab$

```

QED